

# **Domain Name System**

ECE 50863 – Computer Network Systems

# DNS: Mapping between Name and Address

Routers send packets to Internet hosts based on IP addresses (e.g., 216.58.192.196)

Why do we need names? (e.g., www.google.com)

How do we efficiently locate resources?

- DNS: name → IP address

www.google.com -> 216.58.192.196

## Challenge

- How do we scale these to the wide area?

# **Why not centralize DNS?**

Single point of failure

Traffic volume

Distant centralized database

Single point of update

Doesn't scale!

# Domain Name System Goals

A wide-area distributed database

Scalability


Decentralized maintenance

Robustness

Global scope

- Names mean the same thing everywhere

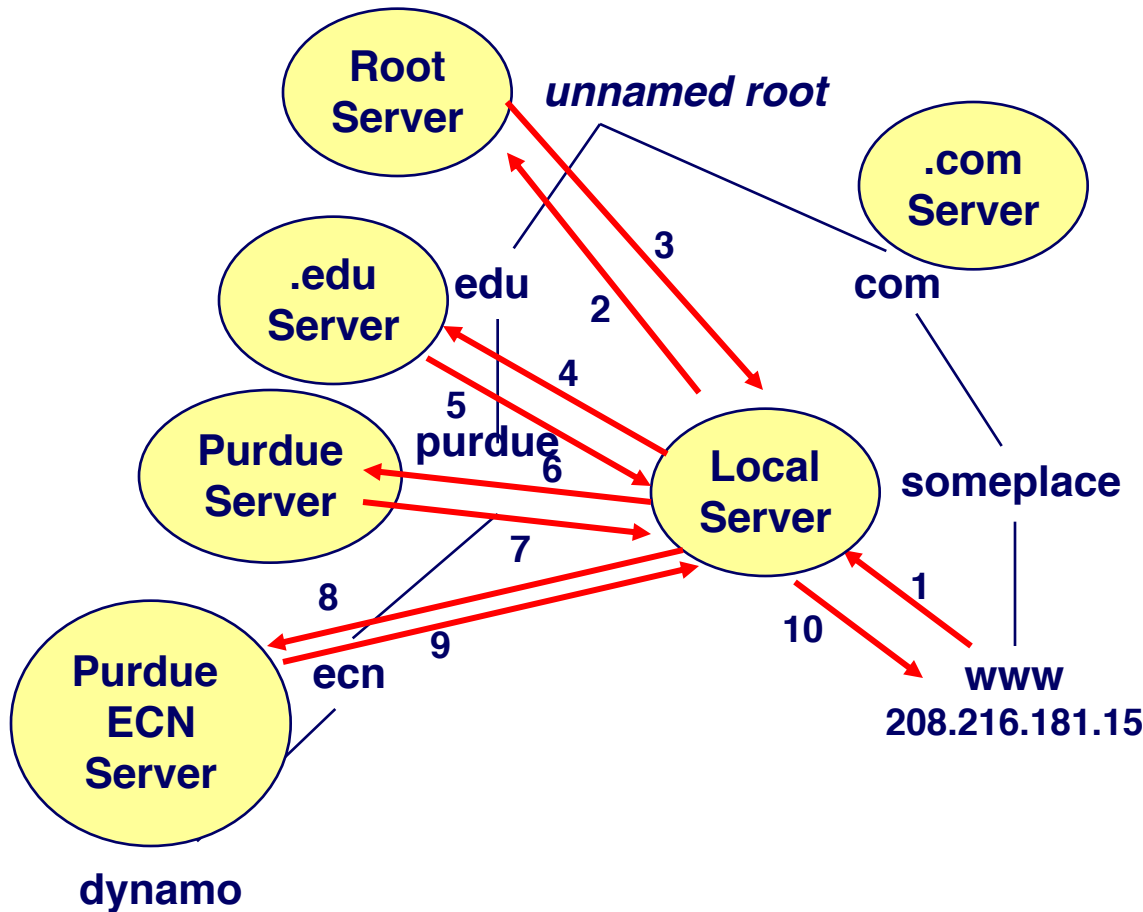
# DNS Design: Hierarchy Definitions



- Hostnames reflect hierarchy
  - E.g., `www.ece.purdue.edu`
- Each node in hierarchy runs one or more “name servers”.
- Leaf level name servers hold name to IP bindings for hosts that end with same suffix
- Name servers at higher levels maintain pointers to name servers of their child zone

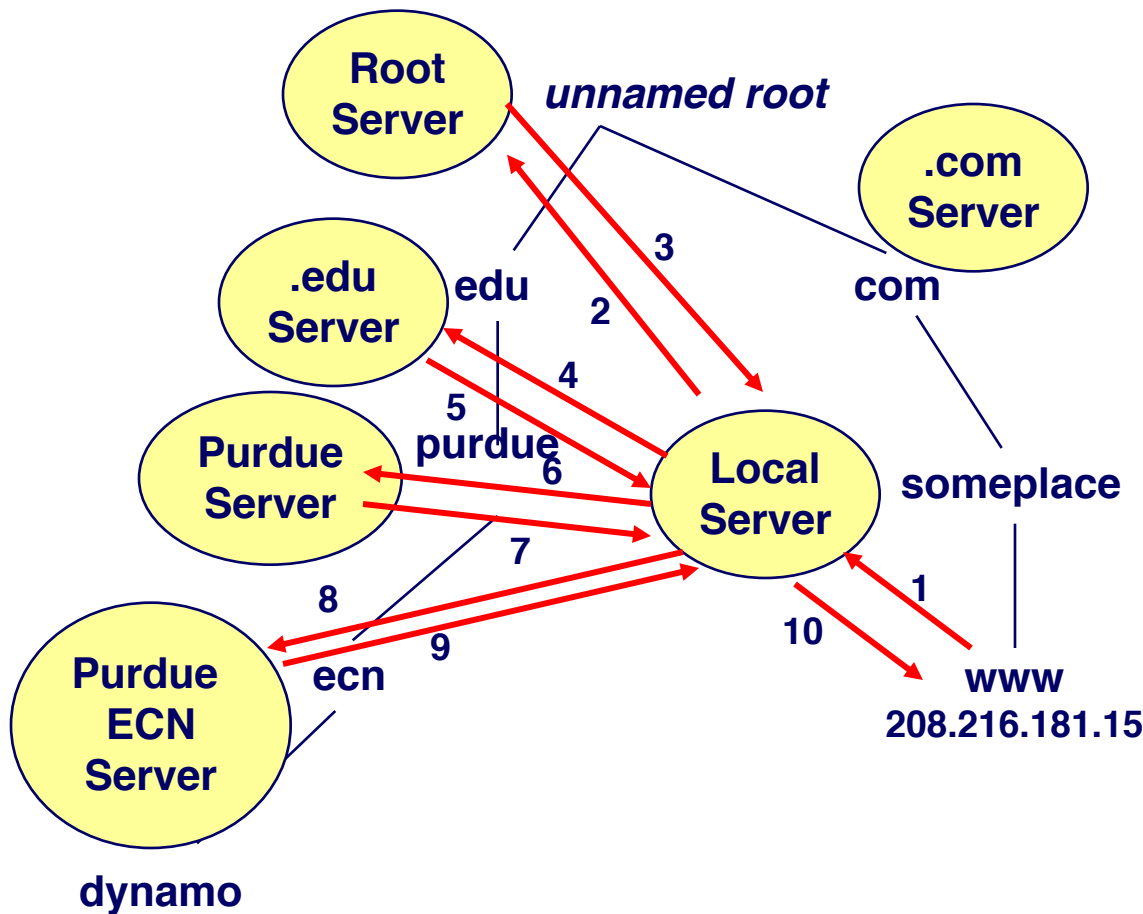
- Purdue ECE’s name server holds bindings for hostnames ending with `.ece.purdue.edu`
- `.edu` name server maintains name servers for `purdue.edu`, `.mit.edu` etc.
- Top level DNS servers maintained by organization called ICANN
  - Heavy replication to ensure robustness.

# Iterative DNS Name Resolution



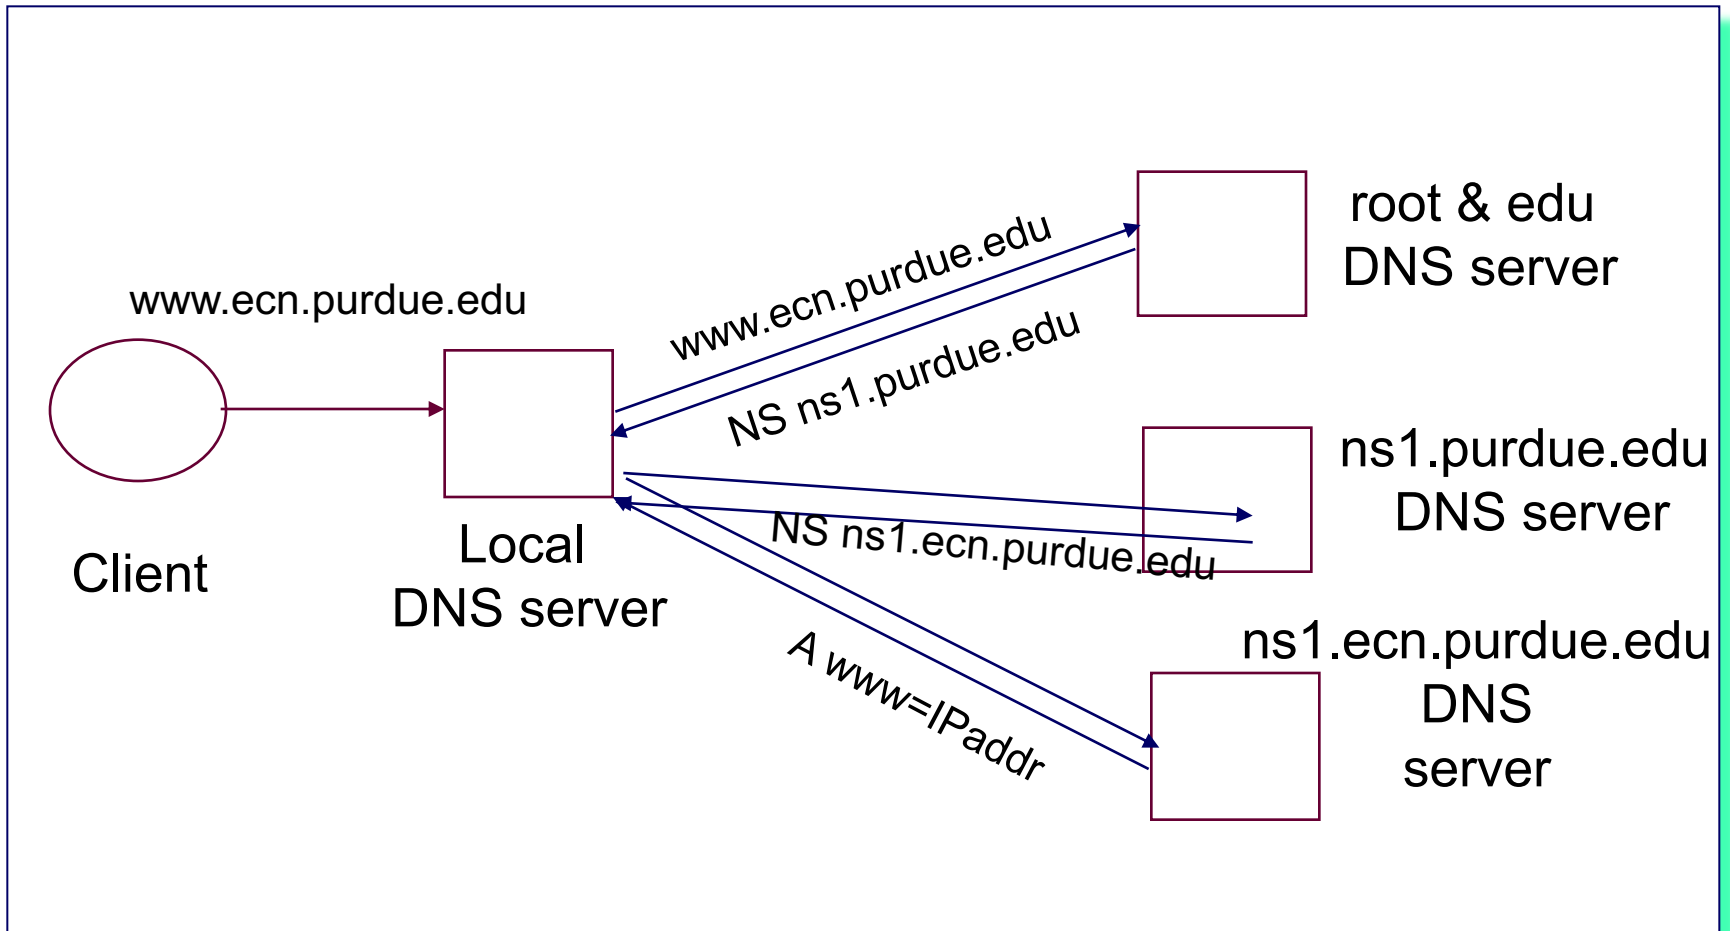
- Local servers do lookup of distant host names for local hosts
- At each step, server returns name of next server down
- Local server directly queries each successive server

# Caching



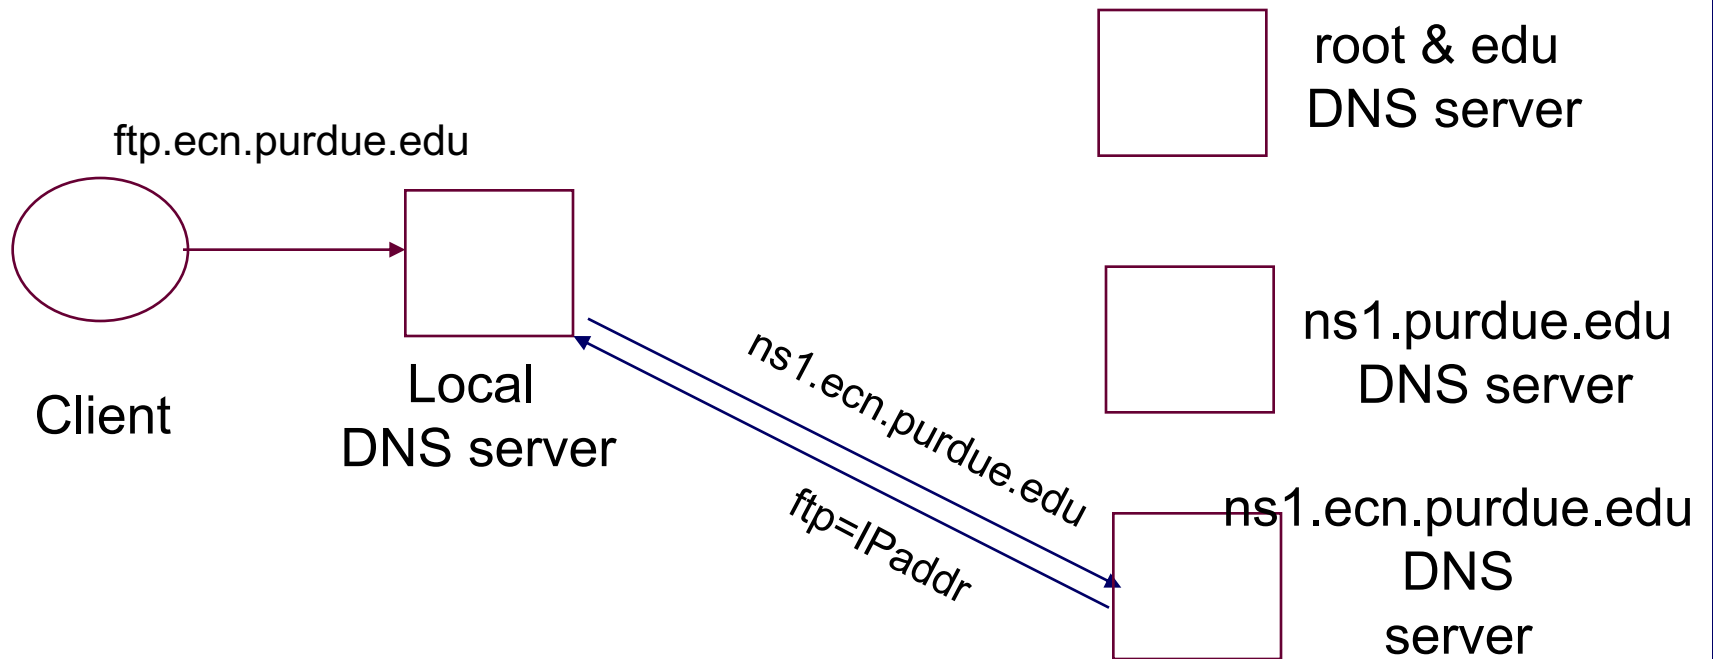
- Local server builds up cache of intermediate translations
- Helps in resolving names  
`xxx.ecn.purdue.edu`,  
`yy.purdue.edu`, and  
`z.edu`
- Cached data periodically times out
  - Lifetime (TTL) of data controlled by owner

# Typical Resolution





# Subsequent Lookup Example



# DNS Records

RR format: (class, name, value, type, ttl)

**DB contains tuples called resource records (RRs)**

- Classes = Internet (IN), Chaosnet (CH), etc.
- Each class defines value associated with type
- Each DNS record associated with its own TTL

## FOR IN class:

### **Type=A**

- name is **hostname**
- value is **IP address**

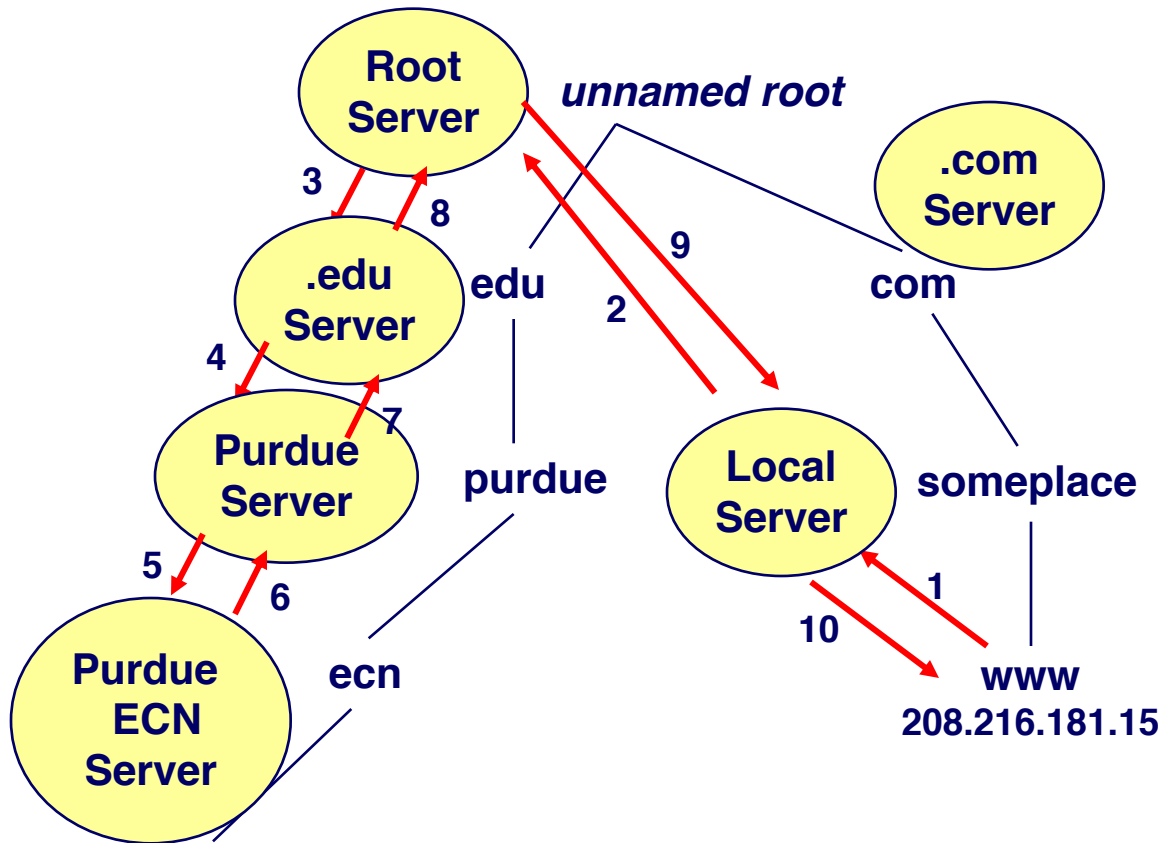
### **Type=NS**

- name is **domain (e.g. foo.com)**
- value is **name of authoritative name server for this domain**

### **Type=CNAME**

- name is **an alias name for some “canonical” (the real) name**
- value is **canonical name**

# Recursive DNS Name Resolution




- Recursively from root server downward
- Results passed up

dynamo  
128.46.200.24

Exploited in security attacks.  
Often disabled, especially at root servers

# Reverse DNS



## Task

- Given IP address, find its name

## Method

- Maintain separate hierarchy based on IP names
- Write 128.46.200.24 as 24.200.46.128.in-addr.arpa

## Managing

- Authority manages IP addresses assigned to it
- E.g., Purdue manages name space 46.128.in-addr.arpa

# Content Delivery Networks and DNS

## Content Delivery Networks (CDNs)

- Replicate content across geographically disperse servers
- Direct users to servers close to them
- First-party CDNs: Google, Netflix etc.
- Third-party CDNs: Akamai, Fastly, CloudFront etc.

## How do CDNs direct clients to close locations?

- Exploit DNS

# CDNs and DNS

## Example redirection process

